

In [2]:

```
import numpy as np
import pandas as pd
import scipy.stats as sp
import seaborn as sns
import matplotlib.pyplot as plt
from pandas.plotting import register_matplotlib_converters
from sklearn.linear_model import LinearRegression as lm
import statsmodels.api as sm
```

In [3]:

```
def difference(dataset, interval=1):
    diff=[]
    for i in range(interval, len(dataset)):
        value=dataset[i]-dataset[i-interval]
        diff.append(value)
    return diff
```

In [4]:

```
def runMyAR1(yin):
    tlen = len(yin)
    y = np.array(yin[2:tlen])
    x = np.array(yin[1:(tlen-1)])
    X = x
    X = sm.add_constant(X)
    regr2 = sm.OLS(y,X)
    model = regr2.fit()
    print(model.summary())
    ypred = model.predict()
    plt.plot((y-ypred))
```

In [5]:

```
# Create Large images!
register_matplotlib_converters()
sns.set_style("darkgrid")
plt.rc("figure", figsize=(14, 8)) # was 16,12
plt.rc("font", size=13)
```

In [6]:

```
dfo = pd.read_csv("Coal Power.csv")
#dfo.columns = ['Month', 'Price']
dfo.head()
```

Out[6]:

	Unnamed: 0	Total consumption : Texas : electric power (total) : quarterly (short tons)
0	2001 Q1	22164839
1	2001 Q2	22952510
2	2001 Q3	25962808
3	2001 Q4	21357650
4	2002 Q1	21917084

In [7]:

```
dfo.rename(columns={"Unnamed: 0": "Period", "Total consumption : Texas : electric power (total) : quarterly (short tons)": "E-Power"})
```

In [8]:

```
dfo.head()
```

Out[8]:

	Period	E-Power
0	2001 Q1	22164839
1	2001 Q2	22952510
2	2001 Q3	25962808
3	2001 Q4	21357650
4	2002 Q1	21917084

In [9]:

```
df = dfo['E-Power']
print(len(df))
```


80

In [10]:

```
df.plot()
```

Out[10]:

<AxesSubplot:>



In [11]:


```
#import chart_studio.plotly as py
import plotly.graph_objs as go
# Offline mode
import plotly.io as pio
#from plotly.offline import init_notebook_mode, iplot
#init_notebook_mode(connected=True)
```

In [12]:

```
#pltobj = go.Scatter(y=dfo['Price'], x=dfo['Month'])
pltobj = go.Scatter(y=dfo['E-Power'])
```

In [13]:

```
fig = go.Figure(data=pltobj)
pio.show(fig)
```



In [14]:

```
sm.tsa.stattools.adfuller(df)
```

Out[14]:

```
(1.9763029096941296,  
 0.9986413567364828,  
 10,  
 69,  
 {'1%': -3.528889992207215,  
  '5%': -2.9044395987933362,  
  '10%': -2.589655654274312},  
 2117.1250184879855)
```

In [15]:

```
sm.tsa.stattools.kpss(df)
```

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\satsmodels\tsa\stattools.py:2018: InterpolationWarning:

The test statistic is outside of the range of p-values available in the look-up table. The actual p-value is smaller than the p-value returned.

Out[15]:


```
(0.8419408149624352,
 0.01,
 4,
 {'10%': 0.347, '5%': 0.463, '2.5%': 0.574, '1%': 0.739})
```

In [16]:

```
# detrend, if required, and plot
dtrend = difference(np.array(df),1)
#dtrend = df
plt.plot(dtrend)
```

Out[16]:

```
[<matplotlib.lines.Line2D at 0x26d80060eb0>]
```



In [17]:

```
sm.tsa.stattools.adfuller(dtrend)
```

Out[17]:

```
(-4.510983475470285,
 0.0001878466316890073,
 5,
 73,
 {'1%': -3.5232835753964475,
 '5%': -2.902030597326081,
 '10%': -2.5883710883843123},
 2087.3684260623113)
```

In [18]:

```
sm.tsa.stattools.kpss(dtrend)
```

Out[18]:

```
(0.4356019919994831,
 0.061809486207119374,
 15,
 {'10%': 0.347, '5%': 0.463, '2.5%': 0.574, '1%': 0.739})
```

In [19]:

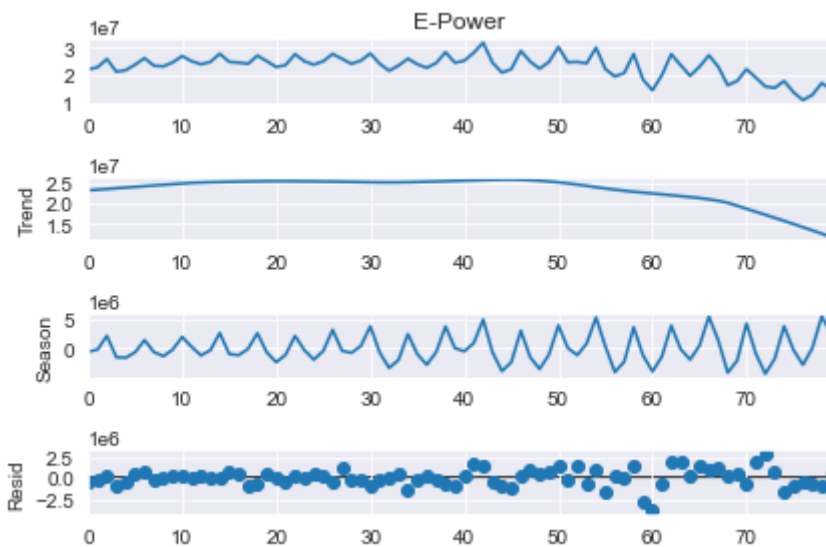
```
from statsmodels.tsa.seasonal import STL
```

In [20]:

```
stl = STL(df, period=12)
#stl = STL(df)
```

In [21]:

```
res = stl.fit()
fig = res.plot()
```



## ARIMA MODEL

In [22]:


```
import statsmodels.api as sm
from statsmodels.tsa.stattools import acf
from statsmodels.tsa.stattools import pacf
```

In [23]:

```
fig = plt.figure(figsize=(14, 8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(dtrend, lags=17, ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(dtrend, lags=17, ax=ax2)
```

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\satsmodels\graphics\tsaplots.py:348: FutureWarning:

The default method 'yw' can produce PACF values outside of the  $[-1,1]$  interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.



In [24]:

```
sm.tsa.stattools.arma_order_select_ic(dtrend)
```

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\satsmodels\tsa\statespace\sarimax.py:978: UserWarning:

Non-invertible starting MA parameters found. Using zeros as starting parameters.

Out[24]:

```
{'bic':
 0      2883.069229  2599.693050  2600.513358
 1      2619.240609  2600.665017  2604.347110
 2      2580.555328  2575.732558  2555.417793
 3      2548.075730  2529.529455  2516.342855
 4      2513.901381  2515.133141  2515.470113,
'bic_min_order': (4, 0)}
```

In [25]:

```
from statsmodels.tsa.arima.model import ARIMA
```



In [26]:

```
my_model = sm.tsa.arima.ARIMA(df,order=(4,1,0),seasonal_order=(4,1,0,12))  
my_model_res = my_model.fit()  
print(my_model_res.summary())
```

## SARIMAX Results

```

=====
=====
Dep. Variable:                               E-Power    No. Observations:
80
Model:                ARIMA(4, 1, 0)x(4, 1, 0, 12)    Log Likelihood
-1065.876
Date:                Mon, 20 Feb 2023    AIC
2149.752
Time:                19:19:09    BIC
2169.594
Sample:                0    HQIC
2157.603

```

```

- 80
Covariance Type:                opg
=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					
-----					
ar.L1	-0.0510	0.118	-0.433	0.665	-0.282
0.180					
ar.L2	-0.0914	0.078	-1.175	0.240	-0.244
0.061					
ar.L3	-0.1802	0.072	-2.507	0.012	-0.321
-0.039					
ar.L4	0.1620	0.071	2.268	0.023	0.022
0.302					
ar.S.L12	-0.2379	0.085	-2.800	0.005	-0.404
-0.071					
ar.S.L24	-0.1300	0.094	-1.386	0.166	-0.314
0.054					
ar.S.L36	-0.0123	0.104	-0.118	0.906	-0.217
0.192					
ar.S.L48	-0.0026	0.077	-0.034	0.973	-0.153
0.148					
sigma2	3.843e+12	8.44e-15	4.55e+26	0.000	3.84e+12
3.84e+12					

```

=====
=====
Ljung-Box (L1) (Q):                0.95    Jarque-Bera (JB):
0.54
Prob(Q):                0.33    Prob(JB):
0.76
Heteroskedasticity (H):                3.12    Skew:
-0.16
Prob(H) (two-sided):                0.01    Kurtosis:
2.70
=====
=====

```

## Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients
(complex-step).
[2] Covariance matrix is singular or near-singular, with condition numb
er 9.69e+41. Standard errors may be unstable.

```

In [27]:

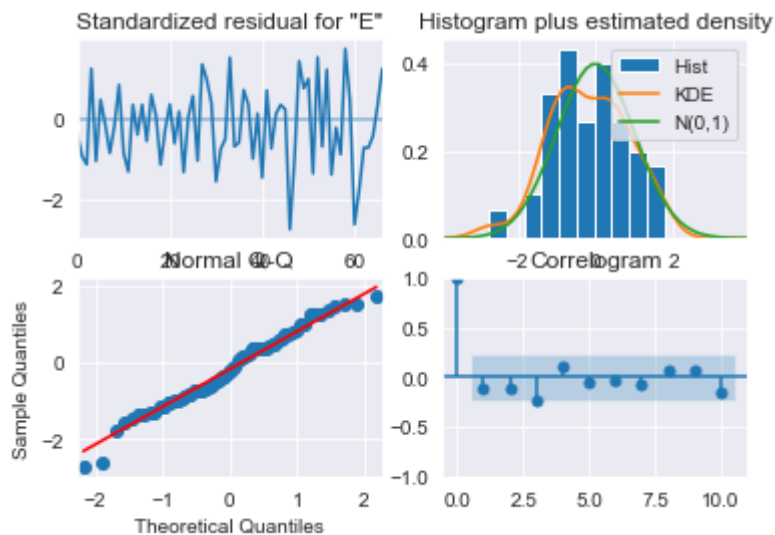
```
type(my_model_res)
```

Out[27]:

```
statsmodels.tsa.arima.model.ARIMAResultsWrapper
```

In [28]:

```
pred = my_model_res.plot_diagnostics()
```



In [29]:

```
tforecast = my_model_res.forecast(24)
tforecast2 = my_model_res.get_forecast(24)
confint = np.array(tforecast2.conf_int())
```

In [30]:

```
type(confint)
```

Out[30]:


```
numpy.ndarray
```

In [31]:

```
plt.plot(df)
plt.plot(tforecast)
```

Out[31]:

```
[<matplotlib.lines.Line2D at 0x26d80542da0>]
```




In [32]:

```
fig,ax = plt.subplots(figsize=(10,5))
ax.plot(df.index, df, label='data')
ax.plot(tforecast2.predicted_mean.index, tforecast2.predicted_mean, label='forecast')
ax.fill_between(tforecast2.predicted_mean.index, confint[:,0], confint[:,1],color='gray')
ax.legend()
```

Out[32]:

```
<matplotlib.legend.Legend at 0x26d80541a50>
```



## AUTO ARIMA MODEL

In [33]:

```
import pmdarima as pm
```

In [34]:

```
myfit = pm.auto_arima(df, start_p=0, start_q=0,
                      max_p=4, max_q=3, m=12,
                      start_P=0, seasonal=True,
                      d=1, D=1, trace=True,
                      error_action='ignore', # don't want to know if an order
                      suppress_warnings=True, # don't want convergence warnin
                      stepwise=True) # set to stepwise
```

Performing stepwise search to minimize aic

```
ARIMA(0,1,0)(0,1,1)[12] : AIC=2156.159, Time=0.24 sec
ARIMA(0,1,0)(0,1,0)[12] : AIC=2151.809, Time=0.04 sec
ARIMA(1,1,0)(1,1,0)[12] : AIC=2152.071, Time=0.19 sec
ARIMA(0,1,1)(0,1,1)[12] : AIC=2151.137, Time=0.23 sec
ARIMA(0,1,1)(0,1,0)[12] : AIC=2151.990, Time=0.07 sec
ARIMA(0,1,1)(1,1,1)[12] : AIC=2152.210, Time=0.57 sec
ARIMA(0,1,1)(0,1,2)[12] : AIC=2151.548, Time=0.69 sec
ARIMA(0,1,1)(1,1,0)[12] : AIC=2151.921, Time=0.17 sec
ARIMA(0,1,1)(1,1,2)[12] : AIC=2153.312, Time=0.99 sec
ARIMA(1,1,1)(0,1,1)[12] : AIC=2152.028, Time=0.62 sec
ARIMA(0,1,2)(0,1,1)[12] : AIC=2151.969, Time=0.30 sec
ARIMA(1,1,0)(0,1,1)[12] : AIC=2151.313, Time=0.17 sec
ARIMA(1,1,2)(0,1,1)[12] : AIC=2153.393, Time=0.56 sec
ARIMA(0,1,1)(0,1,1)[12] intercept : AIC=2152.160, Time=0.26 sec
```

Best model: ARIMA(0,1,1)(0,1,1)[12]

Total fit time: 5.118 seconds

## WINTER-HOLTS MODEL

In [35]:

```
rolling = dfo['E-Power'].rolling(12)
type(rolling)
```

Out[35]:


pandas.core.window.rolling.Rolling

In [36]:

```
mav = rolling.mean()  
plt.plot(dfo['E-Power'])  
plt.plot(mav)
```

Out[36]:

[<matplotlib.lines.Line2D at 0x26d869f5600>]



In [37]:

```
type(mav)
```

Out[37]:


pandas.core.series.Series

In [38]:

```
# Try out the following with various values of 'alpha' and evaluate the results  
ewma = dfo['E-Power'].ewm(alpha=0.5, adjust=False).mean()  
plt.plot(dfo['E-Power'])  
plt.plot(ewma)
```

Out[38]:

[<matplotlib.lines.Line2D at 0x26d86a5aef0>]




In [39]:

```
plt.plot(dfo['E-Power']-ewma)
```

Out[39]:

[<matplotlib.lines.Line2D at 0x26d88abddf0>]



In [40]:

```
from statsmodels.tsa.holtwinters import SimpleExpSmoothing
```

In [41]:

```
ses = SimpleExpSmoothing(df)
```

In [42]:

```
type(ses)
```

Out[42]:

statsmodels.tsa.holtwinters.model.SimpleExpSmoothing

In [43]:

```
result = ses.fit(smoothing_level=0.1, optimized=False)
```

In [44]:

```
result.summary()
```

Out[44]:

SimpleExpSmoothing Model Results

<b>Dep. Variable:</b>	E-Power	<b>No. Observations:</b>	80
<b>Model:</b>	SimpleExpSmoothing	<b>SSE</b>	915218536886908.250
<b>Optimized:</b>	False	<b>AIC</b>	2409.453
<b>Trend:</b>	None	<b>BIC</b>	2414.217
<b>Seasonal:</b>	None	<b>AICC</b>	2409.986
<b>Seasonal Periods:</b>	None	<b>Date:</b>	Mon, 20 Feb 2023
<b>Box-Cox:</b>	False	<b>Time:</b>	19:19:26
<b>Box-Cox Coeff.:</b>	None		

	coeff	code	optimized
<b>smoothing_level</b>	0.1000000	alpha	False
<b>initial_level</b>	2.2165e+07	l.0	False

In [45]:


```
mypred = result.predict(start=1, end=100)
```

In [46]:

```
plt.plot(df)
plt.plot(mypred)
```

Out[46]:

[&lt;matplotlib.lines.Line2D at 0x26d8da61030&gt;]





In [47]:

```
result.params
```

Out[47]:


```
{'smoothing_level': 0.1,
'smoothing_trend': None,
'smoothing_seasonal': None,
'damping_trend': nan,
'initial_level': 22164839.0,
'initial_trend': nan,
'initial_seasons': array([], dtype=float64),
'use_boxcox': False,
'lamda': None,
'remove_bias': False}
```

In [48]:

```
plt.plot(df)
plt.plot(result.fittedvalues)
plt.plot(result.forecast(20))
```

Out[48]:

[&lt;matplotlib.lines.Line2D at 0x26d8dacc2e0&gt;]



In [49]:

```
result2 = ses.fit() # optimize the values of alpha
```

```
C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\s
tatsmodels\tsa\holtwinters\model.py:915: ConvergenceWarning:
```

```
Optimization failed to converge. Check mle_retvals.
```

In [50]:

```
plt.plot(df)
plt.plot(result2.fittedvalues)
plt.plot(result2.forecast(20))
```

Out[50]:

[<matplotlib.lines.Line2D at 0x26d8db17370>]



In [51]:

```
result2.params
```

Out[51]:

```
{'smoothing_level': 0.2812790697674419,
 'smoothing_trend': nan,
 'smoothing_seasonal': nan,
 'damping_trend': nan,
 'initial_level': 22164839.0,
 'initial_trend': nan,
 'initial_seasons': array([], dtype=float64),
 'use_boxcox': False,
 'lamda': None,
 'remove_bias': False}
```

In [52]:

```
from statsmodels.tsa.holtwinters import Holt
```

In [53]:

```
model = Holt(df, exponential=True)
result = model.fit()
result.params
```

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\satsmodels\tsa\holtwinters\model.py:915: ConvergenceWarning:

Optimization failed to converge. Check mle\_retvals.

Out[53]:

```
{'smoothing_level': 0.18837330823548065,
'smoothing_trend': 0.18331477657668624,
'smoothing_seasonal': nan,
'damping_trend': nan,
'initial_level': 21714307.6637008,
'initial_trend': 1.007728898869419,
'initial_seasons': array([], dtype=float64),
'use_boxcox': False,
'lamda': None,
'remove_bias': False}
```

In [54]:

```
result.summary()
```

Out[54]:

Holt Model Results

<b>Dep. Variable:</b>	E-Power	<b>No. Observations:</b>	80
<b>Model:</b>	Holt	<b>SSE</b>	766833837996925.625
<b>Optimized:</b>	True	<b>AIC</b>	2399.301
<b>Trend:</b>	Multiplicative	<b>BIC</b>	2408.829
<b>Seasonal:</b>	None	<b>AICC</b>	2400.452
<b>Seasonal Periods:</b>	None	<b>Date:</b>	Mon, 20 Feb 2023
<b>Box-Cox:</b>	False	<b>Time:</b>	19:19:29
<b>Box-Cox Coeff.:</b>	None		

	coeff	code	optimized
<b>smoothing_level</b>	0.1883733	alpha	True
<b>smoothing_trend</b>	0.1833148	beta	True
<b>initial_level</b>	2.1714e+07	l.0	True
<b>initial_trend</b>	1.0077289	b.0	True

In [55]:

```
plt.plot(df)
plt.plot(result.fittedvalues)
plt.plot(result.forecast(20))
```

Out[55]:

```
[<matplotlib.lines.Line2D at 0x26d8021ff70>]
```



In [56]:

```
from statsmodels.tsa.holtwinters import ExponentialSmoothing
```

In [57]:

```
model = ExponentialSmoothing(df, trend='mul', seasonal='mul', seasonal_periods=12)
```

In [58]:

```
result3 = model.fit()
result3.params
```

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\tsa\holtwinters\model.py:915: ConvergenceWarning:

Optimization failed to converge. Check mle\_retvals.

Out[58]:


```
{'smoothing_level': 0.5353571428571429,
'smoothing_trend': 0.025493197278911566,
'smoothing_seasonal': 0.0001,
'damping_trend': nan,
'initial_level': 24082946.07222221,
'initial_trend': 1.0055785269054447,
'initial_seasons': array([0.91797085, 0.97842296, 1.12838805, 0.991717
93, 0.98135232,
1.01447578, 1.12906098, 0.97248173, 0.88556166, 0.93586909,
1.08454598, 0.98015267])),
'use_boxcox': False,
'lamda': None,
'remove_bias': False}
```

In [59]:

```
plt.plot(result3.level)  
plt.plot(result3.trend)  
plt.plot(result3.season)
```

Out[59]:

[<matplotlib.lines.Line2D at 0x26d802a5d50>]



In [60]:

```
result3.summary()
```

Out[60]:

ExponentialSmoothing Model Results

Dep. Variable:	E-Power	No. Observations:	80
Model:	ExponentialSmoothing	SSE	266642232073999.000
Optimized:	True	AIC	2338.792
Trend:	Multiplicative	BIC	2376.905
Seasonal:	Multiplicative	AICC	2350.005
Seasonal Periods:	12	Date:	Mon, 20 Feb 2023
Box-Cox:	False	Time:	19:19:31
Box-Cox Coeff.:	None		

	coeff	code	optimized
smoothing_level	0.5353571	alpha	True
smoothing_trend	0.0254932	beta	True
smoothing_seasonal	0.0001	gamma	True
initial_level	2.4083e+07	l.0	True
initial_trend	1.0055785	b.0	True
initial_seasons.0	0.9179709	s.0	True
initial_seasons.1	0.9784230	s.1	True
initial_seasons.2	1.1283880	s.2	True
initial_seasons.3	0.9917179	s.3	True
initial_seasons.4	0.9813523	s.4	True
initial_seasons.5	1.0144758	s.5	True
initial_seasons.6	1.1290610	s.6	True
initial_seasons.7	0.9724817	s.7	True
initial_seasons.8	0.8855617	s.8	True
initial_seasons.9	0.9358691	s.9	True
initial_seasons.10	1.0845460	s.10	True
initial_seasons.11	0.9801527	s.11	True

In [61]:

```
plt.plot(df)
plt.plot(result3.fittedvalues)
plt.plot(result3.forecast(24))
```

Out[61]:

```
[<matplotlib.lines.Line2D at 0x26d802782e0>]
```



In [62]:

```
model2 = ExponentialSmoothing(df, trend='add', seasonal='mul', seasonal_periods=12)
```

In [63]:

```
result4 = model2.fit()
result4.params
```

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\s  
tatsmodels\tsa\holtwinters\model.py:915: ConvergenceWarning:

Optimization failed to converge. Check mle\_retvals.

Out[63]:

```
{'smoothing_level': 0.5353571428571429,
'smoothing_trend': 0.025493197278911566,
'smoothing_seasonal': 0.0001,
'damping_trend': nan,
'initial_level': 24082946.07222221,
'initial_trend': 134347.36262626387,
'initial_seasons': array([0.91797085, 0.97842296, 1.12838805, 0.991717
93, 0.98135232,
1.01447578, 1.12906098, 0.97248173, 0.88556166, 0.93586909,
1.08454598, 0.98015267]),
'use_boxcox': False,
'lamda': None,
'remove_bias': False}
```

In [64]:

```
result4.summary()
```

Out[64]:

ExponentialSmoothing Model Results

<b>Dep. Variable:</b>	E-Power	<b>No. Observations:</b>	80
<b>Model:</b>	ExponentialSmoothing	<b>SSE</b>	266092287431320.188
<b>Optimized:</b>	True	<b>AIC</b>	2338.627
<b>Trend:</b>	Additive	<b>BIC</b>	2376.739
<b>Seasonal:</b>	Multiplicative	<b>AICC</b>	2349.840
<b>Seasonal Periods:</b>	12	<b>Date:</b>	Mon, 20 Feb 2023
<b>Box-Cox:</b>	False	<b>Time:</b>	19:19:34
<b>Box-Cox Coeff.:</b>	None		

	coeff	code	optimized
smoothing_level	0.5353571	alpha	True
smoothing_trend	0.0254932	beta	True
smoothing_seasonal	0.0001	gamma	True
initial_level	2.4083e+07	l.0	True
initial_trend	1.3435e+05	b.0	True
initial_seasons.0	0.9179709	s.0	True
initial_seasons.1	0.9784230	s.1	True
initial_seasons.2	1.1283880	s.2	True
initial_seasons.3	0.9917179	s.3	True
initial_seasons.4	0.9813523	s.4	True
initial_seasons.5	1.0144758	s.5	True
initial_seasons.6	1.1290610	s.6	True
initial_seasons.7	0.9724817	s.7	True
initial_seasons.8	0.8855617	s.8	True
initial_seasons.9	0.9358691	s.9	True
initial_seasons.10	1.0845460	s.10	True
initial_seasons.11	0.9801527	s.11	True




In [65]:

```
plt.plot(df)
plt.plot(result4.fittedvalues)
plt.plot(result4.forecast(20))
```

Out[65]:

[<matplotlib.lines.Line2D at 0x26d8eb41e40>]



In [66]:

```
plt.plot(result4.fittedvalues)
```

Out[66]:

[<matplotlib.lines.Line2D at 0x26d8eb985e0>]

